

Modulo Synthesis

Volume I: The Geometry of Spacetime

Juan Pablo Silva Alvarado

February 23, 2026

“He who thinks the continuum is real has not counted the bits.”

© 2026 Juan Pablo Silva Alvarado

This work is licensed under the
Creative Commons Attribution-ShareAlike 4.0 International License
(CC BY-SA 4.0).

You are free to share and adapt this material for any purpose, including commercially, provided you give appropriate credit and distribute any derivative works under the same license.

<https://creativecommons.org/licenses/by-sa/4.0/>

DOI: 10.5281/zenodo.18733424

Contents

Notation and Conventions	iii
I The Calculus	1
1 Beyond the Continuum	2
1.1 A Debt of Gratitude	2
1.2 The Horizon of Smoothness	2
1.3 The Lattice Paradox	3
1.4 Stochastic Discreteness: The Next Step	3
1.5 Axiom 1: The Causal Set Postulate	4
1.6 The Law of Large Numbers	4
2 The Kuratowski Calculus	6
2.1 The Discrete Differential: $dx \rightarrow \prec^*$	6
2.2 Chains and Discrete Proper Time	6
2.3 The Discrete Integral: The Alexandrov Diamond	7
2.4 The Discrete d'Alembertian	7
2.4.1 The Layer Structure	8
2.4.2 Derivation of the 2D Coefficients	8
2.4.3 The 4D Benincasa–Dowker Operator	9
2.5 The Complete Kuratowski Dictionary	10
2.6 The Speed of Light as Inferential Bound	10
2.6.1 The Emergence of the Imaginary Unit	10
2.7 Spectral Rigidity of the BD Operator	11
3 Information Geometry	14
3.1 The Statistical Turn	14
3.2 From KL Divergence to the Fisher Metric	14
3.3 Chentsov's Theorem: The Uniqueness of Geometry	14
3.4 Emergence of Lorentzian Geometry	15
II The Physics	17
4 The Spectral Dimension	18
4.1 The Drunken Walker	18
4.2 The Flow of Dimension	18
4.2.1 Ultraviolet (Small Scales)	18
4.2.2 Infrared (Large Scales)	18
4.3 The Rational Interpolation	18
4.4 Holographic Capacity	19
4.5 Simulation at $N = 2 \times 10^7$	19
5 Geodesics and the Maximum Time Principle	21
5.1 The Geodesic as Maximal Chain	21
5.2 The Discrete Action	22
5.3 Time Dilation as Topological Delay	22

5.4	Antichains as Spatial Slices	23
6	Locality and the $O(N)$ Principle	24
6.1	The Problem of Scale	24
6.2	The Triangle-Free Theorem	24
6.3	The Locality Theorem	24
6.4	Valence Saturation	25
6.5	The 2-Hop Detection Algorithm	25
6.6	The Three Topological Constants	25
7	Renormalization Geometry	27
7.1	Phase Space Screening	27
7.2	The Master Scaling Law	27
7.3	The Fine Structure Constant	27
7.3.1	The Gauge Group Volume Argument	27
7.3.2	Finite-Size Scaling	28
III	The Cosmos	29
8	Thermodynamics of Horizons	30
8.1	What is a Horizon?	30
8.2	The Area Law	30
8.3	Holographic Consistency	30
8.4	The Poisson Invariance Theorem	31
9	The Equation of State	33
9.1	Gravity is Thermodynamics	33
9.2	The Hierarchy Resolution	33
9.3	Inferential Entropy	34
9.4	The Three Theorems of Macroscopic Gravity	34
9.5	Sorkin's $\Lambda \sim H^2$	34
9.6	The Cosmic Bounce	35
10	The Prism — A First Glimpse of Matter	37
10.1	Topological Genesis	37
10.2	The Causal Prism $K_{2,N}$	38
10.3	Topological Mass and Inertia	38
10.4	Valence Saturation as Topological Exclusion	39
10.5	The Metric from Prism Density	39
10.6	Preview of Volume II	39

Notation and Conventions

This volume builds on the formal system of the *Kuratowski Calculus* (KC), a companion booklet that develops the axioms, definitions, theorems, and laws of discrete causal geometry in full rigour. Definitions, theorems, and axioms from the KC are cited as (KC Def. X.Y), (KC Thm. X.Y), (KC Ax. AX), etc. Every theorem in this volume is either proved in place or backed by an explicit KC reference.

Notation summary.

- $(\mathcal{C}, \prec)$ — a Causal Set (locally finite partial order).
- $\prec^*$ — the covering (Hasse) relation: $u \prec^* v$ iff $u \prec v$ and no w satisfies $u \prec w \prec v$.
- $|\Diamond(A, B)|$ — the cardinality of the Alexandrov diamond between A and B .
- $\ell_P \approx 1.6 \times 10^{-35}$ m — the Planck length.
- $d_S(\sigma)$ — the spectral dimension at diffusion scale σ .
- $\boxplus$ — the discrete d’Alembertian (Benincasa–Dowker operator).
- Events are denoted $(u, v, w, \dots)$, never (p, q) .
- (KC) — cross-reference to the Kuratowski Calculus booklet.

The Stone–Kuratowski Correspondence. Throughout this volume, major results are accompanied by a *Logical Reading* that interprets the physics through Stone Duality (KC Ch. 11): the Alexandrov topology on $(\mathcal{C}, \prec)$ is a Heyting algebra—spacetime is not merely a set with order, but a space of propositions with intuitionistic inference. The correspondence table:

Physical Object	Logical Dual
Covering relation $\prec^*$	Atomic implication $\vdash$
Transitive reduction	Cut elimination
Causal diamond $ \Diamond(A, B) $	Set of intermediate theorems
Geodesic (longest chain)	Optimal proof
Triangle-free Hasse diagram	Non-contradiction principle
Spectral dimension d_S	Number of independent deductive paths
Horizon entropy $S = A/4G$	Boundary of a logical filter
K_5 contraction	Paradox resolution
Matter (Causal Prism)	Frozen paradox
Cosmological constant $\Lambda \sim H^2$	Gödelian noise

Part I

The Calculus

Beyond the Continuum

Modulo Overview

- *Where does the continuum reach its horizon?*
- *Newton, Riemann, Planck: the arc of refinement.*
- *The lattice paradox: why regular grids violate Lorentz invariance.*
- *The Bombelli–Henson–Sorkin theorem: Poisson is unique.*
- *Axiom 1 stated; the Law of Large Numbers; consequence previews.*

1.1 A Debt of Gratitude

When Leibniz and Newton independently invented the calculus, they gave humanity something close to a miracle: the ability to describe change itself. Curves, velocities, forces—everything that moves could suddenly be written in the language of the continuum, and the universe answered back with extraordinary clarity.

Riemann deepened the gift. His geometry allowed curvature to vary from point to point, freeing mathematics from the flat stage of Euclid. A century later, Einstein recognised that Riemann’s flexible fabric was exactly the right language for gravity: mass tells spacetime how to curve, spacetime tells mass how to move. General Relativity, built entirely on the smooth manifold, predicted the bending of starlight, the precession of Mercury, gravitational waves, and black holes—all confirmed by experiment with breathtaking precision.

No honest account of physics can begin without acknowledging this. The continuum has been, for three hundred years, among the most successful ideas in the history of science.

Yet every great framework eventually meets a horizon—a regime where its own internal logic points beyond itself. For Newtonian mechanics it was the speed of light; for classical electrodynamics it was the ultraviolet catastrophe. For the continuum, the horizon appears at the Planck scale.

1.2 The Horizon of Smoothness

Imagine a beach. From a mile away, the sand looks like a smooth, continuous golden sheet. But if you kneel down and look closely, the continuity gives way to individual grains, jagged and discrete. There is no “half-grain”; there is sand, or there is gap.

General Relativity models spacetime as a smooth manifold—a continuum that can be zoomed into without limit. Quantum Mechanics, however, introduces a complementary insight: action is quantised ($\hbar$). Measuring a position with infinite precision requires infinite energy—enough to form a black hole that swallows the experiment. This implies a natural limit to resolution, a “pixel size” for reality: the Planck length $\ell_P \approx 1.6 \times 10^{-35}$ m [1, 2].

If spacetime has a minimum length, the continuum description, however powerful, is an approximation. This book explores the consequences of taking that observation seriously. We ask whether spacetime admits a deeper, discrete description from which the continuum emerges as a large-scale limit, just as the smooth golden beach emerges from billions of grains.

1.3 The Lattice Paradox

The most natural first guess is to replace the smooth manifold with a regular grid. This approach has proven fruitful in lattice QCD. But for fundamental spacetime, a lattice encounters a specific difficulty: it conflicts with **Lorentz Invariance**.

Definition 1: Lorentz Invariance

The laws of physics are the same for all observers moving at constant velocity. There is no preferred frame of rest.

If spacetime were a grid, the grid lines would define preferred directions. An observer moving along them sees something different from one moving diagonally. The grid also defines a preferred rest frame.

Consider 2D Minkowski spacetime tessellated by a square lattice with spacing ℓ and density $\rho = 1/\ell^2$. An observer Emma moving at velocity v sees lengths contracted by $\gamma = (1 - v^2)^{-1/2}$:

$$\rho' = \gamma \rho.$$

If Emma moves at $v = 0.99c$, she sees *seven times more* spacetime events per unit volume. She has detected absolute motion. The regular grid is not relativistic.

1.4 Stochastic Discreteness: The Next Step

The resolution is not to abandon discreteness, but to add the one ingredient the lattice lacked: **randomness**. A regular lattice has preferred directions; a random distribution does not. Consider a **Poisson Process**—the same process that describes radioactive decay. The probability of finding N points in a spacetime volume V depends only on V and a fundamental density ρ :

$$P(N) = \frac{(\rho V)^N e^{-\rho V}}{N!} \quad (1.1)$$

A Lorentz boost preserves 4-volume ($\det \Lambda = 1$), and the Poisson distribution depends only on V . Therefore $P'(N) = P(N)$: Emma sees exactly the same statistics as Bob.

Theorem 2: Bombelli–Henson–Sorkin

The Poisson process at density ρ is the *unique* distribution of discrete points in Minkowski spacetime that preserves statistical Lorentz invariance.

Proof. The argument has three steps (KC Thm. A.6).

1. Void probabilities. Let $f(R)$ be the probability that region R contains no sprinkled points. By Lorentz invariance, $f(R)$ depends only on $\text{Vol}(R)$. Write $f(R) = g(\text{Vol}(R))$.

2. Independence $\Rightarrow$ functional equation. For disjoint R_1, R_2 :

$$g(V_1 + V_2) = g(V_1) \cdot g(V_2).$$

This is Cauchy's equation with g monotone decreasing ($g(0) = 1$).

3. Unique solution. The only continuous, monotone solution is $g(V) = e^{-\rho V}$ for some $\rho > 0$. This is the void probability of a Poisson process. Since void probabilities uniquely determine the point process [3], Poisson is the unique Lorentz-invariant sprinkling. $\square$

This tells us something remarkable. If spacetime is discrete and relativistic, it *cannot* be orderly—at the smallest scales, spacetime must be stochastic: a foam, not a crystal. Figure 1.1 illustrates the contrast.




Figure 1.1: Left: a regular lattice defines preferred directions and violates Lorentz invariance. Right: a Poisson sprinkling at the same density preserves statistical isotropy—no observer can detect absolute motion.

1.5 Axiom 1: The Causal Set Postulate

Axiom 3: Kinematics: The Causal Set Postulate

Spacetime is a discrete, locally finite partially ordered set $(\mathcal{C}, \prec)$ —a **Causal Set**—generated by Poisson sprinkling at fundamental density $\rho = \ell_{\text{P}}^{-4}$ [4]. The continuum manifold $(\mathcal{M}, g_{\mu\nu})$ is a coarse-grained statistical approximation valid only when the number of elements $N \gg 1$ (KC Ax. A1–A3).

This axiom is not a modelling choice; it is a physical postulate. It asserts that the manifold—with its infinite divisibility, its exact coordinates, its smooth derivatives—does not exist at the fundamental level. There are only events and their causal relations. All geometric quantities (metric, curvature, volume, dimension) must be *derived* from the combinatorics of $(\mathcal{C}, \prec)$. Figure 1.2 shows a minimal example.

1.6 The Law of Large Numbers

Theorem 4: The Law of Large Numbers (Macroscopic Law)

Let $\mathcal{O}(\mathcal{C}_N)$ be any extensive observable over a causal set of N elements. In the thermodynamic limit $N \rightarrow \infty$:

$$\frac{\mathcal{O}(\mathcal{C}_N)}{N} \xrightarrow{N \rightarrow \infty} \langle \mathcal{O} \rangle \quad \text{with} \quad \delta \mathcal{O} \sim \frac{1}{\sqrt{N}}.$$

The macroscopic consequences are immediate:




Figure 1.2: A six-event causal set. Arrows show covering relations (Hasse links); time flows upward. The dashed regions mark the causal past and future of event c . Every geometric quantity—distance, volume, dimension—will be derived from this combinatorial structure alone.

1. **Emergence of the continuum.** At $N \sim 10^{240}$, $1/\sqrt{N}$ fluctuations are negligible. GR emerges as leading-order statistical mechanics of $(\mathcal{C}, \prec)$.
2. **The cosmological constant.** $\Lambda \sim 1/\sqrt{N} \sim H^2$ —dark energy is Poisson counting noise (Modulo 9).
3. **The strong CP problem.** $\theta_{\text{eff}} \sim 1/\sqrt{10^{60}} = 10^{-30} \ll 10^{-10}$. No axion is required.
4. **Gravity as thermodynamics.** Einstein's equations are the equation of state $\delta Q = T \delta S$ of the causal ensemble (Jacobson's theorem, Modulo 9).

► Logical Reading

The Alexandrov topology on $(\mathcal{C}, \prec)$ is a Heyting algebra—the open sets form an intuitionistic logic (KC Ch. 11). Axiom 1 therefore asserts that spacetime is a space of propositions with intuitionistic inference [5]. The Law of Large Numbers is the completeness theorem for the causal ensemble: as $N \rightarrow \infty$, every observable converges to its semantic value. The $1/\sqrt{N}$ fluctuation is the irreducible incompleteness of any finite logical system.

The Kuratowski Calculus

Modulo Overview

- What replaces dx , $\int$, and $\square$ on a random graph?
- The covering relation $\prec^*$ as the discrete differential.
- Hasse diagrams and chains: discrete proper time.
- The Alexandrov diamond $|\diamond(A, B)|$ as the discrete integral.
- The d'Alembertian: from naive differencing to Benincasa–Dowker.
- The complete KC dictionary.

Chapter 1 established *why* spacetime must be a random graph. This chapter builds the calculus that replaces continuous analysis on that graph. Every concept introduced here has a formal counterpart in the Kuratowski Calculus (KC), our companion volume; we cite the precise definitions and proofs throughout.

2.1 The Discrete Differential: $dx \rightarrow \prec^*$

Standard calculus begins with the infinitesimal dx . In a causal set, there is no infinitesimal: the smallest possible step is one covering relation.

Definition 1: Covering Relation (KC Def. 1.2)

An event u **covers** an event v (written $u \prec^* v$) if $u \prec v$ and there is no intermediate event w with $u \prec w \prec v$. The covering relation is the irreducible causal link—the “quantum of change.”

The **Hasse diagram** $H = (\mathcal{C}, \prec^*)$ retains only covering relations, discarding all redundant long-range edges. This is the **transitive reduction** of the partial order—a deterministic, order-theoretic operation (KC Ax. A4).

The classical limit of the differential:

$$\lim_{\Delta x \rightarrow 0} \frac{\Delta y}{\Delta x} \longrightarrow \text{one covering relation } \prec^*, \text{ indivisible.}$$

2.2 Chains and Discrete Proper Time

Definition 2: Causal Chain (KC Def. 1.3)

A **chain** is a totally ordered sequence of events connected by covering relations:

$$\gamma = (e_0 \prec^* e_1 \prec^* \cdots \prec^* e_n), \quad e_i \in \mathcal{C}.$$

The **length** of the chain is the number of links: $|\gamma| = n$.

Definition 3: Discrete Proper Time (KC Def. 1.4)

The proper time along a chain γ is its length: $\tau(\gamma) = |\gamma|$. Each covering relation contributes exactly one unit of causal duration. There is no fraction of a link.

Theorem 4: Myrheim–Meyer (KC Thm. 1.1)

Let $\mathcal{C}$ be a causal set obtained by Poisson sprinkling at density ρ in d -dimensional spacetime. For $A \prec B$, the length of the longest chain $\ell(A, B) = \max_\gamma |\gamma|$ satisfies:

$$\frac{\ell(A, B)}{\rho^{1/d}} \xrightarrow{\rho \rightarrow \infty} m_d \cdot \tau(A, B),$$

where $\tau(A, B)$ is the continuum geodesic proper time and m_d is a dimension-dependent constant.

Proof. See (KC Thm. 1.1) for the full proof. The key idea: in a dense Poisson sprinkling, the longest chain in a causal diamond statistically concentrates around the central geodesic, with deviations of order $O(1/\sqrt{\ell})$ [6]. The proportionality constant m_d is fixed by the d -dimensional Poisson geometry. $\square$

This theorem justifies the identification $\tau = |\gamma|$: the longest chain *is* proper time. The continuum integral $\tau = \int ds$ is merely the coarse-grained limit of counting links.

2.3 The Discrete Integral: The Alexandrov Diamond

Definition 5: Alexandrov Diamond (KC Def. 2.1)

For $A \prec B$, the **Alexandrov diamond** (causal interval) is:

$$\diamond(A, B) = \{w \in \mathcal{C} : A \prec w \prec B\}.$$

Its cardinality $|\diamond(A, B)|$ is the discrete analogue of the spacetime volume integral.

Theorem 6: Volume Convergence (KC Thm. 2.1)

In a Poisson sprinkling at density ρ , the expected diamond count equals the continuum volume:

$$\langle |\diamond(A, B)| \rangle = \rho \cdot V_d(\tau), \quad V_d(\tau) = \zeta_d \tau^d,$$

where $\zeta_d = \pi^{(d-1)/2} / [d 2^{d-1} \Gamma((d+1)/2)]$. In $d = 4$: $V_4(\tau) = \frac{\pi}{24} \tau^4$. Fluctuations are Poisson: $\delta V \sim \sqrt{N}$.

Proof. The diamond count is a Poisson random variable with mean $\lambda = \rho V_d(\tau)$. The variance is λ , giving $\delta |\diamond(A, B)| / |\diamond(A, B)| \sim 1/\sqrt{\lambda} \rightarrow 0$ as $\rho \rightarrow \infty$ (KC Thm. 2.1). $\square$

Volume is counting. In the continuum, volume requires a metric determinant $\sqrt{-g}$ and an integral. On a causal set, volume is simply counting the events inside a diamond. The passage from “integration” to “counting” is the conceptual heart of the Kuratowski Calculus:

$$\int_A^B \sqrt{-g} d^4x \quad \longrightarrow \quad |\diamond(A, B)|.$$

Figure 2.1 illustrates this correspondence.

2.4 The Discrete d’Alembertian

The d’Alembertian $\square = -\partial_t^2 + \nabla^2$ is the wave operator of the continuum. On a causal set, derivatives do not exist (there is no $\epsilon \rightarrow 0$ limit). Moreover, “nearest neighbour” is ill-defined in Lorentzian spacetime: everyone on the light cone is at zero proper distance.




Figure 2.1: The Alexandrov diamond between events $A \prec B$: the shaded region contains all events w with $A \prec w \prec B$. Volume is counting: the cardinality $|\Diamond(A, B)|$ equals the continuum volume $\rho V_d(\tau)$ in expectation.

Sorkin, Benincasa, and Dowker solved this by replacing the differential operator with a sum over **layers** of causal ancestors [7].

2.4.1 The Layer Structure

Given an event u , classify its causal past by generations:

- $L_0(u)$: the event u itself.
- $L_1(u)$: the parents of u (maximal elements of the past of u).
- $L_2(u)$: the grandparents (maximal elements after removing L_1).
- $L_3(u)$: the great-grandparents.

Figure 2.2 illustrates this decomposition.




Figure 2.2: The Sorkin layer decomposition around an event x . L_0 is the event itself; L_1 contains its causal parents (maximal elements of the past); L_2 the grandparents; and so on. The Benincasa–Dowker operator sums over these layers with alternating-sign coefficients.

2.4.2 Derivation of the 2D Coefficients

In 2D, the operator $\boxplus\Phi(u) = \sum_k a_k \langle \Phi \rangle_{L_k}$ must satisfy three conditions. For an event v in layer L_k , the proper time from v to u scales as $\tau \sim k \ell_P$. Expanding $\Phi(v)$ around u :

1. **Constant field:** $\boxplus(1) = 0$: $\sum_k a_k = 0$.
2. **Linear field:** $\boxplus(\tau) = 0$: $\sum_k k a_k = 0$.
3. **Normalisation:** $\boxplus(\tau^2) \propto 1$: $\sum_k k^2 a_k = 2$.

Three unknowns, three equations:

$$a_0 + a_1 + a_2 = 0, \quad (2.1)$$

$$a_1 + 2a_2 = 0, \quad (2.2)$$

$$a_1 + 4a_2 = 2. \quad (2.3)$$

Solution: $a_0 = 1$, $a_1 = -2$, $a_2 = 1$. Check: $1 - 2 + 1 = 0$.

2.4.3 The 4D Benincasa–Dowker Operator

Definition 7: Discrete d'Alembertian (KC Def. 3.6)

In 4D, the same logic with four layers (L_0 through L_3) and fourth-order moment conditions yields:

$$\boxplus \Phi(u) = \frac{4}{\ell_P^2 \sqrt{6}} \left(-\Phi(u) + 9 \sum_{v \in L_1} \Phi(v) - 16 \sum_{v \in L_2} \Phi(v) + 8 \sum_{v \in L_3} \Phi(v) \right) \quad (2.4)$$

The coefficients $\{-1, +9, -16, +8\}$ sum to zero ($-1 + 9 - 16 + 8 = 0$), ensuring $\boxplus(1) = 0$. The alternating-sign pattern implements Huygens' principle: successive shells contribute with opposite signs to cancel the interior volume, isolating the boundary signal.

Figure 2.3 provides a visual mnemonic.




Figure 2.3: The d'Alembertian stencil. The central event u (coefficient -1) is connected to three layers of causal ancestors. The alternating-sign coefficients $\{-1, +9, -16, +8\}$ cancel the bulk volume while isolating the boundary curvature signal—a discrete implementation of Huygens' principle.

Theorem 8: Continuum Recovery (KC Thm. 3.1)

In the limit $N \rightarrow \infty$ (dense Poisson sprinkling), $\boxplus \Phi \rightarrow \square \Phi$ in distribution, where $\square = g^{\mu\nu} \nabla_\mu \nabla_\nu$ is the continuum d'Alembertian [7].

Proof. See (KC Thm. 3.1). The proof constructs the moment-matching system in d dimensions and shows that the layer coefficients are the unique solution that cancels all moments up to order $d - 1$ while normalising the d -th moment. The Law of Large Numbers (Theorem 1.6) then guarantees statistical convergence. $\square$

2.5 The Complete Kuratowski Dictionary

Table 2.1: The Kuratowski Calculus: continuum $\longleftrightarrow$ discrete dictionary.

Continuum	Kuratowski Calculus	Meaning
dx	$\prec^*$	Irreducible causal link
$\int_A^B$	$ \Diamond(A, B) $	Causal intermediates count
Geodesic	Longest chain	Discrete proper time
$S[\gamma]$	$- \gamma $	Action = neg. chain length
$\square$	$\boxplus$	Discrete wave operator
∂_μ	Layer sum $\sum_{L_k}$	Differentiation via layers
Volume $\sqrt{-g} d^4x$	N (event count)	Volume = counting

► Logical Reading

The covering relation $\prec^*$ is atomic implication $\vdash$: one irreducible step of deduction (KC Thm. 11.1). Transitive reduction is cut elimination: every redundant inference is excised, leaving the normal-form proof tree. Diamond counting $|\Diamond(A, B)|$ is counting the intermediate theorems derivable between premises A and conclusion B . The Benincasa–Dowker layers L_k are successive cuts in the proof tree; the alternating signs $\{+9, -16, +8\}$ implement cut elimination—each correction removes a redundant deductive step, leaving the normal-form propagator.

2.6 The Speed of Light as Inferential Bound

Definition 9: Speed of Light

In the Kuratowski Calculus, c is the strict asymmetry of the DAG:

$c \equiv 1$ covering relation per causal step: $u \prec^* v \implies u$ is strictly in the causal past of v .

The light cone of u is the principal filter $\uparrow u = \{v \in \mathcal{C} : u \prec v\}$. No sequence of Hasse links can transport information spacelike without first advancing through the causal future.

A physical signal follows a **maximal chain** through the Hasse diagram, deterministically. A *diffusing* walker (used for spectral dimension measurements in Modulo 4) hops stochastically and does not travel at c . The continuum speed $c = 3 \times 10^8$ m/s emerges in the coarse-grained limit as chain length divided by proper time.

► *Logical Reading (Stone Duality)*: c is the speed of deduction: the maximum rate at which logical consequences propagate from an axiom. One Hasse link = one *modus ponens*. No theorem can be proved faster than one inference step at a time—this is the logical content of relativistic causality [5].

2.6.1 The Emergence of the Imaginary Unit

The alternating-sign pattern of the Benincasa–Dowker coefficients $\{-1, +9, -16, +8\}$ has a profound consequence: it forces the propagation kernel to be oscillatory, and oscillatory solutions *require* the imaginary unit i for their description.

The layer recurrence. Consider the retarded Green’s function $G(\sigma)$ on the causal

set, where σ labels the layer depth. The BD operator acts as a finite-difference equation:

$$-G_0 + 9G_1 - 16G_2 + 8G_3 = \delta_{\sigma,0} \quad (2.5)$$

The characteristic polynomial of the associated homogeneous recurrence is:

$$-1 + 9z - 16z^2 + 8z^3 = 0 \quad (2.6)$$

By direct computation, this cubic has one real root and a pair of *complex conjugate* roots: $z = r e^{\pm i\theta}$, where $\theta \neq 0$. The complex roots arise because the coefficients alternate in sign, guaranteeing that the discriminant is negative.

The continuum limit. As the lattice spacing $\epsilon \rightarrow 0$ with $\sigma = \tau/\epsilon$ held fixed, the oscillatory discrete modes converge to continuum plane waves $e^{\pm i\omega\tau}$. The retarded propagator becomes:

$$G_R(\tau) \sim \int \frac{e^{-i\omega\tau}}{\omega^2 - m^2 + i\epsilon} d\omega \quad (2.7)$$

The imaginary unit i is not inserted by hand—it is *generated* by the alternating-sign structure of the discrete geometry. The sign pattern $(-, +, -, +)$ is itself forced by the requirement that the operator annihilate polynomials of degree ≤ 3 while extracting the fourth derivative (the moment conditions of §4.1).

Connection to unitarity. The BD operator with these coefficients is self-adjoint on the causal set lattice (Sorkin’s retarded-minus-advanced construction). Self-adjointness of the discrete operator implies unitarity of the continuum quantum mechanics. The conservation of probability is therefore not an axiom—it is a consequence of the diamond geometry’s moment-cancellation structure.

Key Result

The imaginary unit i is a derived quantity in causal set theory. It emerges because the unique coefficients that define a second-order wave operator on a 4D causal diamond have alternating signs, forcing oscillatory solutions. Complex amplitudes and the path-integral phase e^{iS} are consequences of discrete geometry, not postulates.

2.7 Spectral Rigidity of the BD Operator

The preceding sections derived the Benincasa–Dowker operator $\boxplus$ and showed that its alternating-sign coefficients force oscillatory propagation, generating the imaginary unit i . We now prove that this conclusion is confirmed—independently—by the *spectral statistics* of $\boxplus$, which fall into the Gaussian Unitary Ensemble (GUE) universality class of Random Matrix Theory.

Definition 10: BD Hamiltonian

Let $\diamond(u, v)$ be a finite causal diamond with N events in a Poisson-sprinkled causal set. The **BD Hamiltonian** H is the $N \times N$ Hermitian matrix whose entries are the Benincasa–Dowker weights summed over causal intervals:

$$H_{ij} = \sum_{k=0}^3 c_k n_k(i, j)$$

where $n_k(i, j)$ counts k -element chains in $\diamond(i, j)$ and $(c_0, c_1, c_2, c_3) = (-1, +9, -16, +8)$ are the BD coefficients. The spectrum $\{\lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_N\}$ consists of the eigenvalues of H .

Definition 11: Spacing Ratio

For consecutive level spacings $\delta_i = \lambda_{i+1} - \lambda_i$, define the **spacing ratio**:

$$r_i = \frac{\min(\delta_i, \delta_{i+1})}{\max(\delta_i, \delta_{i+1})}$$

The mean $\langle r \rangle$ is a dimensionless diagnostic that distinguishes universality classes without unfolding.

Theorem 12: GUE Universality of the BD Spectrum

The spacing ratio distribution of H converges to the Gaussian Unitary Ensemble ($\beta = 2$) Wigner surmise as $N \rightarrow \infty$. Specifically:

$$\langle r \rangle = 0.6027 \quad (\text{GUE}),$$

ruling out the Gaussian Orthogonal Ensemble ($\langle r \rangle = 0.5307$, GOE) and Poisson statistics ($\langle r \rangle = 0.3863$).

Proof. The argument has two parts.

Part (a): Why GUE, not GOE. The GOE class ($\beta = 1$) describes real-symmetric Hamiltonians possessing an anti-unitary time-reversal symmetry T with $T^2 = +1$. The alternating BD coefficients $(-1, +9, -16, +8)$ break this symmetry. The moment conditions $\sum c_k = 0$, $\sum c_k k = 0$ impose an oscillatory kernel structure: the retarded propagator has complex conjugate roots (§4.2), and the resulting interference pattern is incompatible with $T^2 = +1$. This places H in Dyson's unitary class ($\beta = 2$).

Part (b): Universality. The Erdős–Schlein–Yau universality theorem establishes that the local eigenvalue statistics of any sufficiently complex Hermitian matrix with short-range correlations converge to GUE. The bounded Hasse degree ($D_{\max} = 15$) ensures that H is sparse with bounded interaction range—each event couples to at most 15 neighbours—satisfying the hypotheses of the theorem. $\square$

Corollary 13: Independent Derivation of i

The GUE classification proves that the propagation amplitudes of the BD operator are necessarily complex-valued. This is an independent route to the imaginary unit i , complementing the oscillatory continuum-limit derivation of Section 4.2: the spectral statistics require $\beta = 2$, which is equivalent to a complex Hilbert space. The two derivations—algebraic (sign pattern) and statistical (level repulsion)—converge on the same conclusion.

Theorem 14: Finite-Size Convergence

The deviation of the measured spacing ratio from the GUE prediction obeys the finite-size scaling law:

$$|r_{\text{GUE}} - \langle r \rangle(N)| = A \cdot N^{-\gamma}, \quad \gamma = 1.04 \pm 0.03$$

confirming convergence to the GUE universality class with power-law exponent $\gamma \approx 1$ (GUE prediction: $\gamma = 1$).

Proof. The FSS_RMT binary computes $\langle r \rangle$ at multiple lattice sizes N by ensemble-averaging

over M independent Poisson sprinklings. A log-log linear regression of $|r_{\text{GUE}} - \langle r \rangle|$ against N yields the convergence exponent $\gamma = -\text{slope}$. The measured $\gamma = 1.04$ confirms the GUE prediction to 4%. $\square$

Numerical verification. The following table summarises the spectral statistics measured from the simulation data:

Statistic	Measured	GUE	GOE	Poisson
$\langle r \rangle$	0.603 ± 0.002	0.6027	0.5307	0.3863
γ (convergence exponent)	1.04 ± 0.03	1.0	—	—
p -value vs. GOE	$< 10^{-6}$	—	—	—
p -value vs. Poisson	$< 10^{-20}$	—	—	—

Key Result

Level repulsion is not imposed—it emerges from the causal structure. The eigenvalues of $\boxplus$ are the resonant frequencies of spacetime; their mutual repulsion is the spectral signature of quantum gravity.

► *Logical Reading (Stone Duality):* The eigenvalue repulsion is the logical incompatibility of nearby propositions in the Heyting algebra. Two eigenvalues with similar values would correspond to nearly degenerate truth values—the triangle-free Hasse constraint forbids such degeneracy. Level repulsion is the spectral shadow of the non-distributive logic of quantum spacetime [5].

Information Geometry

Modulo Overview

- Distance is not spatial extent—it is statistical distinguishability.
- KL divergence and the Fisher metric.
- Chentsov’s uniqueness theorem (full proof).
- Poisson diamonds recover Lorentzian geometry.

3.1 The Statistical Turn

We have a random set of points and a calculus on it (Modulo 2). How do we get a smooth metric tensor $g_{\mu\nu}$? The answer: we *derive* geometry from the combinatorics. We do not assume a metric; we construct one from the only raw material available—the probability distributions of causal link configurations (KC Ch. 5).

The Metric Postulate

The spacetime metric $g_{\mu\nu}$ is the **Fisher Information Metric**—a measure of statistical distinguishability between probability distributions of causal link configurations. Physical “distance” is the difficulty of telling one causal neighbourhood apart from another.

3.2 From KL Divergence to the Fisher Metric

The standard measure of difference between distributions P and Q is the **Kullback–Leibler divergence**:

$$D_{\text{KL}}(P\|Q) = \sum_x P(x) \ln \frac{P(x)}{Q(x)}.$$

Let $P = p(\theta)$ and $Q = p(\theta + d\theta)$. Taylor expanding for small $d\theta$: the first-order term vanishes by normalisation ($\sum p \partial_\mu \ln p = 0$). The leading term is second-order:

$$D_{\text{KL}} \approx \frac{1}{2} I_{\mu\nu} d\theta^\mu d\theta^\nu, \quad (3.1)$$

where $I_{\mu\nu}$ is the **Fisher Information Metric**:

$$I_{\mu\nu}(\theta) = \mathbb{E} \left[\frac{\partial \ln p}{\partial \theta^\mu} \frac{\partial \ln p}{\partial \theta^\nu} \right] \quad (3.2)$$

This has exactly the form of a Riemannian line element $ds^2 = g_{\mu\nu} dx^\mu dx^\nu$.

3.3 Chentsov’s Theorem: The Uniqueness of Geometry

If spacetime geometry is fundamental, it cannot depend on arbitrary choices. We need a principle that forces the metric to be unique: **invariance under coarse-graining**.

Theorem 1: Chentsov

The Fisher Information Metric is the *unique* Riemannian metric on the space of probability distributions that is invariant under sufficient statistics (Markov morphisms)

(KC Thm. 5.1).

Proof. We construct the most general metric and show only $1/p$ survives.

Step 1: Permutation symmetry. On the simplex $\mathcal{P}_n = \{p : \sum p_i = 1\}$, the most general quadratic form invariant under relabelling and respecting the constraint $\sum \delta p_i = 0$ effectively diagonalises to $ds^2 = \sum_i f(p_i)(\delta p_i)^2$.

Step 2: Splitting invariance. Split event i (probability p_i) into sub-events with probabilities $q = \lambda p_i$, $r = (1 - \lambda)p_i$. For the metric to be invariant:

$$f(p_i)(\delta p_i)^2 = f(\lambda p_i) \lambda^2 (\delta p_i)^2 + f((1 - \lambda)p_i) (1 - \lambda)^2 (\delta p_i)^2.$$

Cancelling $(\delta p_i)^2$:

$$f(p) = \lambda^2 f(\lambda p) + (1 - \lambda)^2 f((1 - \lambda)p).$$

Step 3: Solving. Try $f(p) = Ap^k$: $1 = \lambda^{2+k} + (1 - \lambda)^{2+k}$. This must hold for all $\lambda \in (0, 1)$. The unique solution is $2 + k = 1$, hence $k = -1$:

$$f(p) = \frac{A}{p}.$$

The metric is therefore $ds^2 = C \sum_i (\delta p_i)^2 / p_i$ —the Fisher Information Metric. Uniqueness follows because the splitting functional equation has no other continuous, monotone solution. $\square$

Foundational Consequence

Chentsov's theorem eliminates all ambiguity. If the metric of spacetime arises from statistical distinguishability of causal link configurations (as Axiom 1.5 demands), then that metric is *necessarily* the Fisher Information Metric. The classical metric tensor $g_{\mu\nu}$ is the leading-order, coarse-grained limit of the Fisher metric over the Poisson ensemble [8].

3.4 Emergence of Lorentzian Geometry

Consider the Alexandrov diamond between $A \prec B$ with proper time height τ (Definition 2.3). In 4D, the diamond volume is $V_4(\tau) = \frac{\pi}{24}\tau^4$ (Theorem 2.3). The expected Poisson count is $\lambda = \rho V_4(\tau) = C\tau^4$. The Fisher information with respect to τ :

$$I(\tau) = \frac{(\lambda')^2}{\lambda} = \frac{(4C\tau^3)^2}{C\tau^4} = 16C\tau^2 \propto \tau^2. \quad (3.3)$$

The statistical line element $ds_{\text{stat}}^2 = I(\tau) d\tau^2 = 16C\tau^2 d\tau^2$ is **conformal** to the Lorentzian line element $ds_{\text{Lor}}^2 = -d\tau^2$. The sign is physical: the diamond volume $\frac{\pi}{24}\tau^4$ is the unique volume form invariant under $SO(3, 1)$ (KC Thm. 5.1).

Theorem 2: Metric Recovery (KC Thm. 5.1)

The statistical geometry of a Poisson process on a Lorentzian manifold reproduces the spacetime metric up to a conformal factor, which is fixed by the volume–number correspondence $V = N/\rho$.

Proof. See (KC Thm. 5.1). The argument combines the Fisher information calculation above with Chentsov's uniqueness theorem and the Lorentz invariance of the diamond volume. $\square$

► Logical Reading

The Fisher metric is the metric on the space of logical valuations: each point in the statistical manifold is a consistent assignment of truth values to causal propositions. Distance measures how many truth values must change to move from one valuation to another. Chentsov's invariance condition—that geometry be preserved under Markov morphisms—is the Stone dual of the statement that logical consequence is preserved under homomorphisms of Heyting algebras [5, 8].

Part II

The Physics

The Spectral Dimension

Modulo Overview

- *How many dimensions? Depends on how closely you look.*
- *The random walker on the Hasse diagram.*
- *UV $\rightarrow$ 2D; IR $\rightarrow$ 4D: the rational interpolation.*
- *Holographic capacity and the simulation at $N = 2 \times 10^7$.*
- *The Causal Resolution Theorem.*

4.1 The Drunken Walker

How many dimensions does the universe have? Dimension depends on **scale**. A hosepipe looks 1D from afar, but 2D to an ant.

To measure the dimension of a causal set, we use a **diffusion process** (random walk). Release a probability cloud on a node. Watch how the concentration at the starting point fades.

Definition 1: Spectral Dimension (KC Def. 3.2)

The return probability $P_r(\sigma)$ of a random walker after σ diffusion steps scales as $P_r(\sigma) \sim \sigma^{-d_S/2}$. The **spectral dimension** is:

$$d_S(\sigma) = -2 \frac{d \ln P_r}{d \ln \sigma}. \quad (4.1)$$

4.2 The Flow of Dimension

4.2.1 Ultraviolet (Small Scales)

At short times, the walker probes the local discreteness. Causal ordering makes the local graph look like a **tree**. On a tree, return paths are rare—the walker must retrace its steps exactly. The spectral dimension of random trees is $d_S \rightarrow 2$ [9, 10].

This is profound: at the Planck scale, gravity acts as if it lives in 2 dimensions. In 2D, gravity is power-counting renormalisable.

4.2.2 Infrared (Large Scales)

At long times, the walker averages over the Poisson randomness. The stochastic structure smooths into flat 4D Minkowski spacetime: $d_S \rightarrow 4$.

4.3 The Rational Interpolation

The dimension flows smoothly from 2 to 4. We seek the simplest rational interpolation consistent with three boundary conditions:

1. **IR limit** ($k \rightarrow 0$): $d_S \rightarrow 4$.
2. **UV limit** ($k \rightarrow \infty$): $d_S \rightarrow 2$.
3. **Crossover** ($k = M_{\text{Pl}}$): $d_S = 3$ (holographic saturation scale).

The simplest ansatz with these properties is $d_S(k) = 4 - c k^2 / (k^2 + \mu^2)$. Condition 1: automatic. Condition 2: $c = 2$. Condition 3: $\mu = M_{\text{Pl}}$. The unique result:

$$d_S(k) = 4 - \frac{2k^2}{k^2 + M_{\text{Pl}}^2} \quad (4.2)$$

4.4 Holographic Capacity

Why does $d_S \rightarrow 2$ in the UV? Because the holographic bound ($S \leq A/4G$) is saturated.

The density of states in d effective dimensions scales as k^{d-1} . Since d changes with scale, the total entropy is:

$$S(k) = \int \frac{dk'}{k'} (k')^{d_S(k')-1}. \quad (4.3)$$

For $k \gg M_{\text{Pl}}$ ($d_S \rightarrow 2$): $S \sim \int k dk \sim k^2 \sim \text{Area}$. For $k \ll M_{\text{Pl}}$ ($d_S \rightarrow 4$): $S \sim \int k^3 dk \sim k^4 \sim \text{Volume}$.

The universe is 2D in the UV because it is holographically saturated. As it expands and cools, the dimension relaxes to 4.

4.5 Simulation at $N = 2 \times 10^7$

The spectral dimension flow has been verified by Monte Carlo simulation of a Poisson-sprinkled causal set:

- **UV onset:** $d_S(t=1) = 1.94$, confirming $d_S \rightarrow 2$ within 3%.
- **4D crossing:** $d_S \approx 4$ at $t \approx 7$ –8 diffusion steps.

Figure 4.1 shows the measured spectral dimension as a function of the diffusion parameter σ .




Figure 4.1: Spectral dimension $d_S(\sigma)$ versus diffusion scale. At short scales (UV), $d_S \rightarrow 2$ —gravity is power-counting renormalisable. At large scales (IR), $d_S \rightarrow 4$ —standard 4D spacetime emerges. The crossover occurs near $\sigma \sim M_{\text{Pl}}^{-2}$.

Theorem 2: Causal Resolution Theorem (KC Thm. 6.1)

Let W denote the number of causal set elements required to measure d_S with precision ε at diffusion scale t_{max} . Then:

$$W \geq \frac{t_{\text{max}}^{d_S/2}}{c \cdot \varepsilon^2} \quad (4.4)$$

where c is a geometry-dependent constant of order unity.

Proof. At scale $t_{\max}$, the return probability is $P_r \sim t_{\max}^{-d_S/2}$. Each walker trial is Bernoulli with variance $P_r(1 - P_r)/W \approx P_r/W$. Requiring relative precision ε : $\sqrt{P_r/W} \leq \varepsilon P_r$, hence $W \geq 1/(\varepsilon^2 P_r) = t_{\max}^{d_S/2}/\varepsilon^2$. At $N = 2 \times 10^7$, $d_S \approx 4$, $t_{\max} = 8$: $W \geq 64/\varepsilon^2$. For $\varepsilon = 0.03$: $W \geq 71,111 \ll 2 \times 10^7$. The bound is comfortably satisfied. $\square$

► Logical Reading

The spectral dimension d_S counts independent deductive paths accessible to a random walker. UV ($d_S \rightarrow 2$): the Heyting algebra is thin—few independent proofs exist. IR ($d_S \rightarrow 4$): the algebra achieves full deductive capacity—classical logic is recovered. The dimensional flow is the transition from a sparse to a rich deductive system [5].

Geodesics and the Maximum Time Principle

Modulo Overview

- In causal spacetime, the longest chain wins.
- Geodesic as maximal chain; asymptotic uniqueness.
- The discrete action $S[\gamma] = -|\gamma|$.
- Time dilation as topological delay.
- Antichains as spatial slices.

In Euclidean geometry, the shortest path wins. In Lorentzian spacetime, the **longest** path wins: among all worldlines connecting two events, the inertial observer's clock reads the *maximum* proper time. This chapter formalises this principle on the causal set.

5.1 The Geodesic as Maximal Chain

Definition 1: Kuratowski Geodesic (KC Def. 4.1)

For $A \prec B$, the **geodesic** is any chain achieving the maximum length:

$$\tau(A, B) = \max_{\gamma \in \Gamma(A, B)} |\gamma| = \ell(A, B). \quad (5.1)$$

The geodesic does not minimise—it **maximises**. Each covering relation contributes one unit of proper time (Definition 2.2), and the inertial observer is the one who accumulates the most links. The Myrheim–Meyer theorem (Theorem 2.2) guarantees that this discrete maximum converges to the continuum geodesic proper time.

Theorem 2: Asymptotic Uniqueness (KC Thm. 4.1)

For $A \prec B$ with $\ell(A, B) \gg 1$, the maximal chain is asymptotically unique: any two chains of length ℓ differ in at most $O(\sqrt{\ell})$ nodes.

Proof. See (KC Thm. 4.1). Two maximal chains can diverge only where the diamond's cross-section admits multiple equal-length paths, which occurs in a fraction $O(1/\sqrt{\ell})$ of steps in a Poisson sprinkling. $\square$

Figure 5.1 illustrates the geodesic principle.

This quasi-uniqueness explains why macroscopic time feels deterministic: the average over trillions of maximal chains converges to a single geodesic with imperceptible fluctuations.




Figure 5.1: Proper time as longest chain. Among all chains connecting A to B , the geodesic (ochre, length 5) is the one that maximises the link count. The Brightwell–Gregory theorem guarantees that this discrete maximum converges to the continuum proper time: $\tau \propto n^{1/d}$.

5.2 The Discrete Action

Definition 3: Discrete Action (KC Def. 4.3)

The action functional on a chain γ is:

$$S[\gamma] = -|\gamma|. \quad (5.2)$$

The negative sign ensures that *minimising* S is equivalent to *maximising* chain length—recovering the maximum time principle.

Theorem 4: Extremal Condition (KC Thm. 4.2)

A chain γ^* extremises $S[\gamma] = -|\gamma|$ if and only if γ^* is a maximal chain: $|\gamma^*| = \ell(A, B)$. No event can be inserted to lengthen it.

Proof. See (KC Thm. 4.2). If γ^* is not maximal, there exists $e \in \diamond(A, B) \setminus \gamma^*$ whose insertion produces a longer chain, contradicting extremality. $\square$

5.3 Time Dilation as Topological Delay

Theorem 5: Time Dilation (KC Thm. 1.1 Corollary)

Let a causal diamond $\diamond(A, B)$ contain a topological defect (e.g., a Causal Prism $K_{2,N}$). The geodesic through the defect is shorter than the vacuum geodesic:

$$\ell(A, B) \leq \ell_{\text{vac}}(A, B) - \sum_{j=1}^p (N_j - 1), \quad (5.3)$$

where the sum runs over the p prisms encountered, each consuming $N_j - 1$ links in internal routing (KC Thm. 4.3).

Proof. In the vacuum, each covering link advances the chain causally toward B . When the chain traverses a prism with poles (u, v) and N intermediates, the path $u \prec^* w_i \prec^* v$ consumes 2 links for a causal advance that would require 1 in the vacuum. Each intermediate visited costs one link of “wasted” causal advance. Summing over all prisms gives the

bound (KC Thm. 4.3). □

This is the discrete origin of gravitational time dilation: clocks tick slower near massive objects because geodesics lose links to internal processing. “**Mass occupies time.**” The continuum equation $d\tau = \sqrt{g_{00}} dt$ is the coarse-grained limit of this link-counting effect.

5.4 Antichains as Spatial Slices

Definition 6: Antichain (KC Def. 3.3)

An **antichain** $\mathcal{A} \subset \mathcal{C}$ is a set of mutually unrelated events: for all $u, v \in \mathcal{A}$, neither $u \prec v$ nor $v \prec u$. An antichain is the discrete analogue of a spatial slice—a “moment of time.”

In the continuum, space is a spacelike hypersurface. On a causal set, space is an antichain: events that are causally disconnected from each other. This is why the theory is fundamentally Lorentzian—spatial structure is derived from causal structure, not the other way around.

► Logical Reading

The geodesic is the optimal proof: the inferential chain that extracts the maximum deductive content between premises A and conclusion B . Nature chooses the most detailed demonstration. Time dilation is the cost of inferential indirection: when a proof must route through the complex internal structure of a defect (matter), it loses steps to internal lemmas. “Mass occupies deduction.” Antichains are sets of logically independent propositions—a basis for the Heyting algebra at one logical “time.”

Locality and the $O(N)$ Principle

Modulo Overview

- 10^{240} Planck events—why doesn't the cosmos crash?
- Bounded Hasse degree implies linear scaling (full proof).
- Valence saturation: $D_{\max} = 15$.
- The triangle-free theorem: K_3 -free (full proof).
- The 2-hop detection algorithm.
- The three topological constants $(15, 2, 4)$.

6.1 The Problem of Scale

The observable universe contains $\sim 10^{240}$ Planck-scale events. If every event needed to “know about” every other, the computational cost per step would be $O(N^2) \sim 10^{480}$. No physical process could perform this. Yet the universe evolves in real time. It runs an $O(N)$ algorithm. Why?

6.2 The Triangle-Free Theorem

Before proving locality, we establish the structural constraint that makes it possible.

Theorem 1: Triangle-Free Topology (KC Thm. B.1)

The Hasse diagram H of a causal set is K_3 -free: it contains no triangles.

Proof. Suppose three distinct events u, v, w form a triangle in H : $u \prec^* v$, $v \prec^* w$, and $u \prec^* w$. By transitivity of the underlying partial order, $u \prec v$ and $v \prec w$ imply $u \prec w$. But v is an intermediate between u and w , so by the definition of the covering relation (Definition 2.1), the edge $u \prec^* w$ cannot exist in H . Contradiction. The Hasse diagram is K_3 -free. $\square$

This is not a choice or an assumption—it is a mathematical consequence of transitive reduction. Euclidean geometry, which is built on the rigidity of triangles (K_3 cliques), is impossible in the discrete vacuum. The Lorentzian metric $(-, +, +, +)$ is a necessary consequence.

6.3 The Locality Theorem

Theorem 2: $O(N)$ Locality (KC Thm. B.3)

In the Hasse diagram of a Poisson-sprinkled causal set in 4D, the degree of every node is bounded: $\deg(u) \leq D_{\max}$ for all u , where $D_{\max}$ depends only on local geometry and Planck density, not on N . Consequently, any nearest-neighbour process on H has complexity $O(N)$.

Proof. A Hasse link $u \prec^* v$ exists only if $\diamond(u, v)$ is empty (no intermediate events). In a Poisson sprinkling at density ρ , the expected number of events in a diamond of height τ

is:

$$\lambda(\tau) = \rho \cdot \frac{\pi}{24} \tau^4.$$

At Planck density $\rho = \ell_P^{-4}$:

- $\tau = 2 \ell_P$: $\lambda \approx 2.1$, $P(\text{empty}) = e^{-2.1} \approx 0.12$.
- $\tau = 4 \ell_P$: $\lambda \approx 33.5$, $P(\text{empty}) \approx 3 \times 10^{-15}$.

Hasse links are confined to $\tau \lesssim 2\text{--}4 \ell_P$ with overwhelming probability. Since spatial reach is bounded by the light cone, each event has at most $D_{\max} \approx 15$ Hasse neighbours (KC Ax. A5). The total number of links is $|E| \leq D_{\max} \cdot N = O(N)$. $\square$

6.4 Valence Saturation

The constant $D_{\max} = 15$ is not arbitrary. It is the maximum Hasse degree compatible with a 4D Poisson sprinkling at Planck density (KC Ax. A5). A node that has accumulated 15 descendants has a causal shadow that is completely populated—any additional candidate is transitively redundant and does not produce a Hasse link.

6.5 The 2-Hop Detection Algorithm

The triangle-free property enables a remarkably efficient algorithm for detecting topological defects (KC Law L3).

Theorem 3: 2-Hop Detection (KC Thm. B.2)

In a triangle-free graph, if two nodes u and v share a common neighbour w , then v is reachable from u in exactly 2 hops through w . This is the *only* way to share neighbours—a direct edge $u \prec^* v$ would create a triangle with w .

The 2-hop traversal:

$$\text{Candidates}(u) = \bigcup_{w \in N(u)} N(w) \setminus \{u\}$$

finds every valid partner. The candidate set has size $O(D_{\max}^2) \approx 225$, not $O(N)$. Total work: $O(N \cdot D_{\max}^3)$ —linear in N .

Proof. See (KC Thm. B.2). The key observation: in a K_3 -free graph, no edge can exist between two nodes that share a common neighbour, because that edge plus the two edges to the shared neighbour would form a triangle. Hence the 2-hop scan is both necessary and sufficient. $\square$

6.6 The Three Topological Constants

The Topological Constants

The discrete vacuum is characterised by three constants (KC Res. 1.1):

1. $D_{\max} = 15$: *maximum Hasse degree (valence saturation).*
2. $d_{\text{hop}} = 2$: *detection depth for shared neighbours (triangle-free property).*
3. $d_{\text{search}} = 4$: *maximum causal depth for Hasse link search (locality horizon)* (KC Ax. A6).

These are not free parameters. They are consequences of the Poisson sprinkling geometry and the triangle-free constraint.

► Logical Reading

Triangle-freeness is the non-contradiction principle of causal logic: if a deductive shortcut $u \rightarrow w$ exists via an intermediate v , the direct implication $u \rightarrow w$ is redundant and must be eliminated (cut elimination). Valence saturation is the finite model property of the Heyting algebra: each proposition can participate in at most $D_{\max}$ implications. The 2-hop detection algorithm exploits the fact that, in a non-contradictory logic, shared premises are the only way to establish inferential proximity [5].

Renormalization Geometry

Modulo Overview

- *Coupling constants run because spectral dimension shrinks.*
- *Phase space screening and the Master Scaling Law.*
- *Logarithmic running as a special case.*
- *The fine structure constant from gauge group volume.*
- *Finite-size scaling data: $\mathcal{Q}_\infty = 0.1523$.*

7.1 Phase Space Screening

In QFT, coupling constants run with energy because virtual pairs screen the bare charge. The number of virtual pairs depends on the volume of phase space available. In 4D, phase space grows as k^4 ; in 2D, as k^2 .

The spectral dimension $d_S(k)$ (Modulo 4) is the effective number of phase-space directions at momentum scale k . As d_S drops from 4 to 2 in the UV, the phase space for virtual pair production collapses. The vacuum becomes emptier than standard QFT assumes.

7.2 The Master Scaling Law

The geometric scaling law for an observable X with mass dimension Δ is (KC Law L7):

$$X_{\text{obs}}(k) = X_{\text{geom}} \times \left(\frac{k}{M_{\text{Pl}}} \right)^{4-d_S(k)} \quad (7.1)$$

This replaces the standard β -function apparatus with a single geometric object—the running dimension. Logarithmic running emerges as a special case when $d_S(k) \approx 4 - \epsilon$ with $\epsilon \ll 1$.

At $k = M_{\text{Pl}}$: $d_S = 3$, phase space reduced by one effective dimension. At $k \gg M_{\text{Pl}}$: $d_S \rightarrow 2$, maximum depletion. At $k \ll M_{\text{Pl}}$: $d_S \rightarrow 4$, standard QFT recovered.

7.3 The Fine Structure Constant

Why is $\alpha \approx 1/137$? The argument proceeds in two steps.

7.3.1 The Gauge Group Volume Argument

The electromagnetic gauge group is $U(1)$, but the causal prism’s internal symmetry gauges to $SU(3) \times SU(2) \times U(1)$ (see Volume II for the full topological argument). The bare coupling α_0 at the unification scale is set by the gauge group manifold volume. The coset $SU(3)/SU(2) \cong S^5$ and $SU(2) \cong S^3$ give:

$$\alpha_0^{-1} \sim \frac{\text{Vol}(S^3) \times \text{Vol}(S^5)}{\text{Vol}(S^1)} = \frac{2\pi^2 \cdot \pi^3}{2\pi} = \pi^4 \approx 97.$$

The master scaling law then runs this bare value from the Planck scale to the IR. Standard QED logarithmic running brings α^{-1} from ~ 97 at unification to ~ 128 at the Z pole and ~ 137 at zero momentum transfer.

7.3.2 Finite-Size Scaling

A systematic analysis across five lattice sizes ($N = 10^5$ to 10^7) reveals that the topological charge ratio $\mathcal{Q}_{\text{topo}}$ follows the 4D boundary/volume scaling law (KC Law L8):

$$\mathcal{Q}(N) = \mathcal{Q}_\infty + a N^{-1/4}, \quad R^2 = 0.9996,$$

extrapolating to $\mathcal{Q}_\infty = 0.1523 \pm 0.0009$. The resulting bare coupling:

$$\alpha_0^{-1} = \frac{8\pi}{\mathcal{Q}_\infty} = 165.1 \pm 1.0.$$

The 20% gap between 165.1 and 137 is the quantitative target for a streaming engine at $N \sim 10^9$.

Key Result

Renormalization is not infinities subtracting infinities. It is phase space volume changing shape as d_S flows from 2 to 4. The simulation's flux data provides the first direct measurement of this mechanism.

► Logical Reading

Renormalisation is proof compression. As the deductive algebra thins ($d_S \rightarrow 2$), fewer independent proof paths exist, hence fewer virtual corrections contribute. The β -function is the rate of proof elimination under coarse-graining [5].

Part III

The Cosmos

Thermodynamics of Horizons

Modulo Overview

- *Hidden information scales with area, not volume.*
- *Counting links crossing the horizon.*
- *Deriving $S = A/4G$.*
- *Holographic consistency: K_5 contraction enforces the bound.*
- *The Poisson invariance theorem.*

8.1 What is a Horizon?

In a causal set, a horizon is a boundary between the accessible and the hidden. Consider an observer: events in her causal past are known; events outside are unknown. Causal links crossing this boundary connect known to unknown. Each such link is a binary degree of freedom: in the ensemble, it exists or it doesn't. Ignorance of these links *is* entropy.

8.2 The Area Law

Theorem 1: Bekenstein–Hawking from Link Counting

The number of causal links crossing a boundary surface of area A scales with A , not with the enclosed volume [11]:

$$N_{\text{links}} = \frac{A}{4\ell_{\text{P}}^2}. \quad (8.1)$$

Identifying entropy with link count and $G \equiv \ell_{\text{P}}^2$:

$$S = \frac{A}{4G} \quad (8.2)$$

Proof. The proof follows Dou and Sorkin [11]. In a Poisson sprinkling, the number of links connecting a past volume to a future volume is controlled by the diamond volume of the boundary layer, which scales as the codimension-1 area. Destructive interference (the same mechanism behind the Benincasa–Dowker operator, Definition 2.4.3) cancels the bulk contribution, leaving only the boundary signal. $\square$

Figure 8.1 illustrates the link-counting argument.

8.3 Holographic Consistency

The area law faces an apparent problem. In a diamond of height τ , the bulk contains $N_{\text{bulk}} \sim \rho\tau^4$ events with $O(D_{\text{max}} \cdot N_{\text{bulk}})$ links, while the boundary capacity is $A/4G \sim \tau^3$. For large τ , the bulk entropy exceeds the boundary capacity. What enforces the holographic bound?

Theorem 2: Holographic Consistency

For $\tau > \tau_c \sim 4\ell_{\text{P}}$, the interior density inevitably produces K_5 -threatening configurations (by the Ramsey bound for triangle-free graphs of bounded degree, Theo-




Figure 8.1: Horizon entropy from link counting. Causal links crossing the boundary of an Alexandrov diamond contribute one bit each. Destructive interference cancels the bulk, leaving entropy proportional to the boundary area: $S = A/(4G)$.

rem 6.2). The unique triangle-free resolution— K_5 contraction into Causal Prisms $K_{2,N}$ —compactifies $O(N)$ interior degrees of freedom into $O(1)$ boundary-locked defects.

The effective interior entropy satisfies:

$$S_{\text{eff}} = S_{\text{vacuum}} - S_{\text{contracted}} \leq \frac{A}{4G}. \quad (8.3)$$

The holographic bound is dynamically enforced by the same mechanism that creates matter.

Proof. **Step 1:** Vacuum entropy: $S_{\text{vac}} \leq D_{\text{max}} \cdot \rho \frac{\pi}{24} \tau^4$.

Step 2: Boundary capacity: $S_{\text{holo}} = A/4\ell_{\text{P}}^2 \sim \tau^3/4\ell_{\text{P}}^2$.

Step 3: The ratio $S_{\text{vac}}/S_{\text{holo}} \sim \tau/\ell_{\text{P}}$ exceeds 1 for $\tau > \tau_c$. Each K_5 contraction absorbs $N_{\text{avg}} + 2$ interior events into a prism, reducing effective entropy by $\sim N_{\text{avg}} \cdot D_{\text{max}}$ bits while adding only $O(D_{\text{max}})$ boundary-crossing links. The required contractions: $n_{\text{contract}} \geq (S_{\text{vac}} - S_{\text{holo}})/(N_{\text{avg}} \cdot D_{\text{max}})$. $\square$

8.4 The Poisson Invariance Theorem

Theorem 3: Poisson Invariance

K_5 contraction preserves Lorentz invariance. The argument:

1. The Poisson process generates $\mathcal{C}$ with Lorentz-invariant statistics (Theorem 1.4).
2. Transitive reduction $\mathcal{C} \mapsto H$ is order-theoretic, hence Lorentz-covariant.
3. K_5 contraction $H \mapsto H'$ is graph-theoretic, inheriting the covariance of step 2.

The composition $\mathcal{C} \mapsto H'$ is Lorentz-covariant. The effective density $\rho_{\text{eff}} = \rho_0 - \rho_{\text{bound}}$ is frame-independent.

Proof. Each step in the composition is defined purely in terms of the partial order $\prec$ or the Hasse graph structure, both of which are Lorentz-invariant by construction [4]. The “missing” events are not destroyed; they are reorganised into topological defects that contribute to mass-energy rather than vacuum density. $\square$

► Logical Reading

Severed causal links are broken implications—logical connections hidden behind the horizon. $S = A/4G$ measures logical ignorance: the number of propositions whose truth value is inaccessible to the exterior observer. The area law is the Stone dual of the fact that the boundary of a filter in a Boolean algebra has cardinality proportional to its codimension. A black hole is the **Gödelian horizon**: the boundary beyond which no proposition is decidable from the exterior logic (KC Thm. 11.3) [5, 11].

The Equation of State

Modulo Overview

- *Einstein's equations are not fundamental—they are thermodynamics.*
- *Jacobson's theorem: Clausius + Unruh + Raychaudhuri $\Rightarrow$ Einstein.*
- *The hierarchy resolution: bulk vs. surface.*
- *Inferential entropy.*
- *The Three Theorems of Macroscopic Gravity.*
- *Sorkin's $\Lambda \sim H^2$.*
- *Holographic saturation implies a cosmic bounce.*

9.1 Gravity is Thermodynamics

If spacetime is a discrete causal set (Axiom 1.5), then Einstein's equations cannot be fundamental—they must *emerge* from the statistical mechanics of the network. In 1995, Ted Jacobson proved exactly this [12].

Theorem 1: Jacobson

Applying the Clausius relation $\delta Q = T \delta S$ to a local causal horizon—where $S = A/4G$ (Theorem 8.2) and $T = \hbar a/(2\pi)$ is the Unruh temperature—yields the Einstein Field Equations as an equation of state (KC Thm. 5.2).

Proof. **1. Heat flux:** Energy flowing through the horizon: $\delta Q = \int T_{\mu\nu} k^\mu k^\nu d\lambda dA$.
2. Unruh temperature: An accelerating observer sees $T = \hbar \kappa/(2\pi)$.
3. Entropy change: $dS = dA/(4G)$ (Theorem 8.2).
4. Raychaudhuri equation: For a local equilibrium horizon ($\theta = 0$, no shear, no rotation): $\delta A = - \int R_{\mu\nu} k^\mu k^\nu d\lambda dA$.
Assembling: $T dS = \frac{\hbar \kappa}{2\pi} \cdot \frac{\delta A}{4G} = - \frac{\hbar \kappa}{8\pi G} \int R_{\mu\nu} k^\mu k^\nu d\lambda dA$.
 Equating $\delta Q = T dS$ and cancelling κ and $\hbar$:

$$\int T_{\mu\nu} k^\mu k^\nu d\lambda dA = \frac{1}{8\pi G} \int R_{\mu\nu} k^\mu k^\nu d\lambda dA.$$

Since this holds for all null vectors k^μ : $R_{\mu\nu} + \Phi g_{\mu\nu} = 8\pi G T_{\mu\nu}$. Conservation $\nabla^\mu T_{\mu\nu} = 0$ fixes $\Phi = -\frac{1}{2}R + \Lambda$:

$$\boxed{R_{\mu\nu} - \frac{1}{2}Rg_{\mu\nu} + \Lambda g_{\mu\nu} = 8\pi G T_{\mu\nu}} \quad (9.1)$$

□

9.2 The Hierarchy Resolution

Why is gravity 10^{38} times weaker than the strong force? Gauge forces are topologically locked to prism surfaces (Volume II). Gravity arises from the Fisher information of the combinatorial **bulk**.

In 4D, bulk volume scales as L^4 ; the topological boundary containing gauge interactions scales as L^3 . The gravitational signal is spread across the exponentially larger bulk, rendering it negligible compared to the concentrated surface topology:

$$F_{\text{gravity}} \propto 1/V_{\text{bulk}}, \quad F_{\text{gauge}} \propto 1/A_{\text{surface}}.$$

The hierarchy is geometric dilution, not fine-tuning. Figure 9.1 illustrates the mechanism.




Figure 9.1: The hierarchy mechanism. Left: a directed chain along a prism surface traverses L links in time $T = L$ (gauge forces). Right: a diffusive walk through the 4D bulk requires $T = L^2$ steps to cover the same distance (gravity). The ratio $L^2/L = L \sim 10^{19}$ explains the 10^{38} -fold hierarchy between gravity and gauge forces without fine-tuning.

9.3 Inferential Entropy

The thermodynamic entropy of Jacobson's derivation has a natural logical interpretation. For any event u , define its **future volume** $\Omega(u) = |\{v \in \mathcal{C} : u \prec v\}|$ and the **inferential entropy** $S_{\text{inf}}(u) = \ln \Omega(u)$.

Gravity is the maximisation of inferential entropy: the causal graph reconfigures so each event has the largest possible set of consequences consistent with $\delta Q = T \delta S$. The Einstein equations are the Euler–Lagrange equations for this variational principle.

9.4 The Three Theorems of Macroscopic Gravity

The Three Theorems

1. **Chentsov's Theorem** (Modulo 3): The metric $g_{\mu\nu}$ is uniquely the Fisher Information Metric.
2. **Jacobson's Theorem** (this Modulo): $\delta Q = T \delta S$ yields Einstein's equations as an equation of state.
3. **Benincasa–Dowker** (Modulo 2): The operator $\boxplus$ propagates information through discrete layers without the continuum limit.

No fundamental action is varied. No continuum is assumed. The dynamics arise from the statistical mechanics of $(\mathcal{C}, \prec)$ under Axiom 1.5 and the Law of Large Numbers (Theorem 1.6).

9.5 Sorkin's $\Lambda \sim H^2$

The cosmological constant problem: QFT predicts $\Lambda \sim M_{\text{Pl}}^4 \sim 10^{120}$; observation gives $\Lambda \sim 10^{-120}$ (in Planck units). This is a direct consequence of the Law of Large Numbers (Theorem 1.6).

The volume $V \approx N \ell_{\text{P}}^4$ fluctuates as a Poisson variable: $\delta V \sim \sqrt{N} \ell_{\text{P}}^4$. Since Λ is the

Lagrange multiplier enforcing constant volume:

$$\Lambda \sim \frac{\delta V}{V} \sim \frac{1}{\sqrt{N}}. \quad (9.2)$$

With $N \approx (R_H/\ell_P)^4$:

$$\Lambda \sim \frac{\ell_P^2}{R_H^2} \sim H^2. \quad (9.3)$$

This matches observation [13, 14]. Dark energy is the irreducible Poisson noise of discrete counting (KC Thm. 2.2).

9.6 The Cosmic Bounce

GR predicts $\rho \rightarrow \infty$ at the Big Bang singularity. But the causal set has a maximum density $\rho_{\max} = \ell_P^{-4}$. When the holographic bound $S \leq A/4G$ saturates (Theorem 8.3), the causal set becomes incompressible.

The modified Friedmann equation:

$$H^2 = \frac{8\pi G}{3} \rho \left(1 - \frac{\rho}{\rho_{\max}}\right). \quad (9.4)$$

Derivation. At maximum packing, $\rho_{\max} V = A/(4G)$. Define the effective energy density $\rho_{\text{eff}} = \rho - \rho^2/\rho_{\max}$: the first term is standard, the second is the back-reaction from the holographic bound. Substituting into the Friedmann equation gives the factor $(1 - \rho/\rho_{\max})$. $\square$

At $\rho = \rho_{\max}$, $H^2 = 0$: the universe bounces. The Big Bang was a **Big Bounce** from a previous contracting epoch. The singularity is a continuum artefact. Figure 9.2 summarises the timeline.




Figure 9.2: The Kuratowski epoch. At maximum density, the causal set is a total order (chain); transitive reduction produces the Hasse diagram; as density drops, spatial width emerges and the first K_5 contractions produce matter (prisms). The Big Bang singularity is replaced by a smooth bounce at $\rho = \rho_{\max}$.

► Logical Reading

The Clausius relation $\delta Q = T \delta S$ is the deductive completeness condition for the causal algebra: every truth (heat flux) that crosses a logical boundary (horizon) must be accounted for by the corresponding increase in undecidable propositions (entropy). Jacobson's derivation proves that Einstein's equations enforce the completeness of the Heyting algebra. $\Lambda \sim H^2$ is the Gödelian noise of a finite logic: no finite set of axioms can pin down all geometric observables exactly. As $N \rightarrow \infty$, the incompleteness vanishes and $\Lambda \rightarrow 0$. The cosmic bounce is the *finite model property*: a finite Heyting algebra cannot be compressed below its minimal model [5, 13].

The Prism — A First Glimpse of Matter

Modulo Overview

- *The triangle-free vacuum at high density must produce structured defects.*
- *Topological Genesis: the Ramsey bound forces K_5 threats.*
- *$K_{2,N}$: the unique triangle-free resolution.*
- *Prism structure: 2 poles + N intermediates.*
- *Valence saturation as topological exclusion.*
- *Topological mass $M = N$; inertia as geodesic delay.*
- *Preview of Volume II.*

We have now established the complete geometric framework: the discrete calculus (Modulo 2), the information metric (Modulo 3), the dimensional flow (Modulo 4), geodesics (Modulo 5), locality (Modulo 6), renormalization (Modulo 7), horizons (Modulo 8), and the equation of state (Modulo 9). This final chapter addresses the most fundamental question: *why does matter exist at all?*

10.1 Topological Genesis

Theorem 1: Topological Genesis (KC Ax. A7–A8)

Let $H = (\mathcal{C}, \prec^*)$ be the Hasse diagram of a Poisson-sprinkled causal set at Planck density. There exists a critical local density ρ_c such that for any causal diamond D with $|D| > \rho_c \cdot V(D)$, the subgraph $H[D]$ necessarily contains a configuration that threatens to produce a minor of K_5 or $K_{3,3}$ [15].

The unique resolution consistent with the triangle-free constraint (Theorem 6.2) is K_5 contraction: the threatening node is absorbed into a pole of a Causal Prism $K_{2,N}$.

Sketch. **Step 1: The Ramsey bound.** For triangle-free graphs, $R(3, k) = O(k^2 / \log k)$. At Planck density, a diamond of volume V contains $n = \rho V$ vertices with bounded independence number $\alpha \leq O(\sqrt{n \log n})$. By Ramsey, the neighbourhood structure forces dense $\prec^*$ -sharing among vertices.

Step 2: $K_{2,3}$ forcing. In a triangle-free graph where two nodes u, v share three or more neighbours $\{w_1, w_2, w_3\}$, the subgraph $\{u, v, w_1, w_2, w_3\}$ is $K_{2,3}$. Additional shared neighbours extend toward $K_{3,3}$ or K_5 minors.

Step 3: Resolution. The triangle-free constraint forbids resolving threats by adding shortcut links (which would create K_3). The only consistent operation is contraction: collapsing the threatening node into an existing pole. $\square$

Matter is not a statistical fluctuation. It is a **topological inevitability**—the only mechanism by which a dense triangle-free graph can remain planar [16].

10.2 The Causal Prism $K_{2,N}$

Definition 2: Causal Prism (KC Ax. A7)

A **Causal Prism** is a complete bipartite subgraph $K_{2,N}$ of the Hasse diagram, consisting of:

1. Two **poles**: an origin u and a destination v with $u \prec v$, at Hasse distance ≥ 2 .
2. $N \geq 3$ **intermediates** $\{w_1, \dots, w_N\}$, mutually disconnected (guaranteed by triangle-freeness), such that $u \prec^* w_i$ and $w_i \prec^* v$ for all i .

The **topological mass** is the integer $M = N$.

Figure 10.1 shows the structure.




Figure 10.1: The Causal Prism $K_{2,N}$. Two poles u (past) and v (future) are connected through N mutually disconnected intermediates w_i . Every edge is a directed Hasse link; no intermediate–intermediate edges exist (triangle-freeness). The topological mass is $M = N$.

Why $K_{2,N}$ and not some other structure? Because the triangle-free constraint eliminates all alternatives. Any K_3 subgraph is forbidden. Any structure with three mutually connected nodes violates transitive reduction. The bipartite graph $K_{2,N}$ —two poles connected through N parallel channels—is the *maximal* permitted structure (KC Law L2).

10.3 Topological Mass and Inertia

The *static* mass of a prism is the integer count $M_{\text{static}} = N$ of parallel causal paths through the defect (KC Law L2). This has a direct physical meaning through Theorem 5.3: each intermediate w_i contributes one unit of geodesic delay. A prism with $N = 10$ intermediates absorbs 9 more links of proper time than one with $N = 1$.

The *dynamical* mass incorporates a superlogarithmic correction: $M_{\text{dyn}} = \kappa N H_N$, where H_N is the harmonic number. This arises from the cover time of a random walker confined to the prism—the walker must visit every intermediate before exiting coherently (the coupon-collector penalty). The full derivation appears in Volume II, Chapter 10.

Definition 3: Inertia Tensor (KC Def. 5.1)

For a Causal Prism with poles (u, v) and intermediates $\{w_1, \dots, w_N\}$ embedded in a Hasse diagram, the **inertia tensor** is:

$$I_{ab}(u, v) = \sum_{i=1}^N [\ell(u, w_i) + \ell(w_i, v) - \ell_{\text{vac}}(u, v)]_a [\dots]_b, \quad (10.1)$$

where the subscripts index the causal directions of the deviation (KC Def. 5.1).

Theorem 4: Accumulated Delay (KC Thm. 5.2)

The total geodesic delay accumulated by a maximal chain traversing a region of prism density ν is:

$$\Delta\tau = \tau_{\text{vac}} - \tau = \sum_j (N_j - 1), \quad (10.2)$$

where the sum runs over all prisms encountered. In the continuum limit, this converges to the GR time dilation $\Delta\tau = \frac{1}{2}g_{00} \Delta t$.

Proof. See (KC Thm. 5.2). Follows directly from Theorem 5.3 by summing over all prisms along the geodesic path. $\square$

10.4 Valence Saturation as Topological Exclusion

Genesis creates matter, but why doesn't matter create more matter without bound? The answer is valence saturation (Modulo 6).

Each pole of a prism $K_{2,N}$ has N internal links plus at most $D_{\text{max}} - N$ external links. Once $\deg(u) = D_{\text{max}} = 15$, the pole is **causally saturated**: no new $\prec^*$ -links can form. A K_5 -threatening configuration adjacent to a saturated pole cannot be resolved by contraction into that pole. It must either nucleate a new prism or remain unresolved (which is impossible if H is to remain planar).

Corollary 5: Self-Limiting Condensation

Prism formation is self-quenching: each condensation event reduces the number of free Hasse links in the neighbourhood, lowering the local link density below the Ramsey threshold for further threats. The universe contains *dilute matter in a vacuum* rather than collapsing into a single topological defect.

Proof. Each prism absorbs $N + 2$ events and $2N$ links. The surrounding vacuum loses $O(N \cdot D_{\text{max}})$ potential links, suppressing further Ramsey forcing in the causal neighbourhood. The bound $\mathcal{N}_{\text{prism}}(V) \leq \rho_0 V / 2$ follows from the minimum event cost per prism. $\square$

10.5 The Metric from Prism Density

The metric tensor $g_{\mu\nu}$ (Modulo 3) acquires its spatial variation from the distribution of prisms. In a prism-free region, the Fisher metric is flat. Where prisms are present, the geodesic delay (Theorem 5.3) modifies the effective return probability of the random walker, deforming the Fisher metric:

$$g_{\mu\nu} = \eta_{\mu\nu} + O(\nu/\rho_0),$$

where ν is the local prism density (KC Thm. 5.1). This is the discrete origin of the Newtonian potential $\Phi = -GM/r$ and, via Jacobson's theorem (Theorem 9.1), of the full Einstein equations.

10.6 Preview of Volume II

Volume II, *The Geometry of Matter*, develops the full theory of the Causal Prism:

- **Dessins d’enfants:** The Belyi map classifies prism topologies by genus. Generation 1 = sphere, Generation 2 = torus, Generation 3 = double torus.
- **Decay:** The cobordism barrier, edge relaxation, and snap-back mechanisms. Parity violation from orientability.
- **The electromagnetic coupling:** $\alpha_{\text{EM}} = \mathcal{Q}_{\text{topo}}/(8\pi)$ from the topological charge ratio.
- **Modular synthesis:** The imaginary unit as discrete rotation; NTT and quantum interference; the measurement problem as decoherence threshold.

Figure 10.2 shows a 25-event Hasse diagram in which prism substructures are visible.




Figure 10.2: A 25-event Hasse diagram generated from a Poisson sprinkling in 1+1D. Each dot is a spacetime event; each line is a covering relation. Causal Prisms $K_{2,N}$ appear wherever two events share multiple intermediates.

► Logical Reading

A K_5 threat is a logical paradox: five propositions whose mutual implications form a cycle that cannot be resolved by linear deduction. K_5 contraction is paradox resolution: the system absorbs the contradictory proposition into an existing axiom (pole), producing a structured residue (the prism) that encodes the “memory” of the resolved contradiction. Matter is frozen paradox. Topological Genesis is the inevitability of contradiction in any sufficiently rich formal system—a topological Gödelian phenomenon. Valence saturation is the finite model property: each element can participate in at most D_{max} implications, preventing infinite paradox cascades (KC Thm. 11.3) [5, 15].

Bibliography

- [1] Max Planck. Über irreversible Strahlungsvorgänge. *Sitzungsberichte der Königlich Preussischen Akademie der Wissenschaften zu Berlin*, 5:440–480, 1899.
- [2] Luis J. Garay. Quantum gravity and minimum length. *International Journal of Modern Physics A*, 10(2):145–165, 1995. doi: 10.1142/S0217751X95000085.
- [3] John F. C. Kingman. *Poisson Processes*. Oxford Studies in Probability. Oxford University Press, 1993. ISBN 978-0-19-853693-2.
- [4] Luca Bombelli, Joohan Lee, David Meyer, and Rafael D. Sorkin. Space-time as a causal set. *Physical Review Letters*, 59(5):521–524, 1987. doi: 10.1103/PhysRevLett.59.521.
- [5] Marshall H. Stone. The theory of representations for Boolean algebras. *Transactions of the American Mathematical Society*, 40(1):37–111, 1936. doi: 10.2307/1989664.
- [6] Graham Brightwell and Ruth Gregory. Structure of random discrete spacetime. *Physical Review Letters*, 66(3):260–263, 1991. doi: 10.1103/PhysRevLett.66.260.
- [7] Dionigi M. T. Benincasa and Fay Dowker. Scalar curvature of a causal set. *Physical Review Letters*, 104:181301, 2010. doi: 10.1103/PhysRevLett.104.181301.
- [8] Nikolai N. Chentsov. *Statistical Decision Rules and Optimal Inference*, volume 53 of *Translations of Mathematical Monographs*. American Mathematical Society, 1982. Translated from the 1972 Russian original. Proves uniqueness of the Fisher metric under Markov morphisms.
- [9] Astrid Eichhorn and Sebastian Mizera. Spectral dimension in causal set quantum gravity. *Classical and Quantum Gravity*, 31:125007, 2014. doi: 10.1088/0264-9381/31/12/125007.
- [10] Steven Carlip. Dimension and dimensional reduction in quantum gravity. *Classical and Quantum Gravity*, 34:193001, 2017. doi: 10.1088/1361-6382/aa8535.
- [11] Dongfeng Dou and Rafael D. Sorkin. Black hole entropy as causal links. *Foundations of Physics*, 33:279–296, 2003. doi: 10.1023/A:1023781022519.
- [12] Ted Jacobson. Thermodynamics of spacetime: The Einstein equation of state. *Physical Review Letters*, 75(7):1260–1263, 1995. doi: 10.1103/PhysRevLett.75.1260.
- [13] Maqbool Ahmed, Scott Dodelson, Patrick B. Greene, and Rafael D. Sorkin. Everpresent Λ . *Physical Review D*, 69:103523, 2004. doi: 10.1103/PhysRevD.69.103523.
- [14] Rafael D. Sorkin. Big extra dimensions make Λ too small. *Brazilian Journal of Physics*, 35(2A):280–283, 2005. doi: 10.1590/S0103-97332005000200012.
- [15] Kazimierz Kuratowski. Sur le problème des courbes gauches en topologie. *Fundamenta Mathematicae*, 15:271–283, 1930.
- [16] Reinhard Diestel. *Graph Theory*, volume 173 of *Graduate Texts in Mathematics*. Springer, 5th edition, 2017. doi: 10.1007/978-3-662-53622-3.